

Bias and Fairness in LLMs

Trustworthy Large Language Models

ELL881 · AIL821



Yatin Nandwani
Research Scientist, IBM

Bias and Fairness in LLMs



Bias and Fairness in LLMs

Business News > Technology > Artificial Intelligence > Meta sued by 26 ex-employees over alleged AI-driven layoff bias

Meta sued by 26 ex-employees over alleged AI-driven layoff bias

The lawsuit, filed in the US District Court in Oakland, California, claims Meta relied on AI-assisted performance metrics —including productivity levels and AI token usage — when selecting employees for job cuts.



Bias and Fairness in LLMs

Business News > Technology > Artificial Intelligence > Meta sued by 26 ex-employees over alleged AI-driven layoff bias

Meta sued by 26 ex-employees over alleged AI-driven layoff bias

The lawsuit, filed in the US District Court in Oakland, California, alleges that Meta used AI-driven metrics—including productivity levels and AI token usage—to

hai.stanford.edu/news/ai-hiring-tools-can-yield-racial-bias-and-systemic-rejection

NEWS

AI Hiring Tools Can Yield Racial Bias and Systemic Rejection



Bias and Fairness in LLMs

Business News > Technology > Artificial Intelligence > Meta sued by 26 ex-employees over alleged AI-driven layoff bias

Meta sued by 26 ex-employees over alleged AI-driven layoff bias

The lawsuit, filed in the US District Court in Oakland, California, alleges that Meta used AI-driven metrics—including productivity levels and AI token usage—to

hai.stanford.edu/news/ai-hiring-tools-can-yield-racial-bias-and-systemic-rejection

NEWS

AI Hiring Tools Can Yield Racial Bias and Systemic Rejection



World ▾ Business ▾ Markets ▾ Sustainability ▾ Legal ▾ Commentary ▾ Technology ▾ Investig

Workday must face California lawsuit over AI bias in job screening tools

By Daniel Wiessner

June 23, 2026 1:42 AM GMT+5:30 · Updated June 23, 2026



Trustworthy Large Language Models



Yatin Nandwani

Bias and F


MEDIA

AI chatbots take heat over left-wing bias: 'No longer be considered neutral'

Business News > Technology > Artificial Intelligence > Meta sued by 26 ex-employees over alleged AI-driven layoff bias

Meta sued by 26 ex-employees over alleged AI-driven layoff bias

The lawsuit, filed in the US District Court in Oakland, California, alleges that Meta used AI-driven metrics—including productivity levels and AI token usage—to make layoff decisions.



World ▾ Business ▾ Markets ▾ Sustainability ▾ Legal ▾ Commentary ▾ Technology ▾ Investic

Workday must face California lawsuit over AI bias in job screening tools

By Daniel Wiessner

June 23, 2026 1:42 AM GMT+5:30 · Updated June 23, 2026



Trustworthy Large Language Models



Yatin Nandwani

Bias and Fairness in LLMs

What is wrong here?

Who is harmed — and how?

Is it bias? Is it unfairness? Are those the same thing?



Bias and Fairness in LLMs

What is
Fairness/Bias?



Bias and Fairness in LLMs


What is
Fairness/Bias?




Why LLMs are
biased?



Bias and Fairness in LLMs



Bias and Fairness in LLMs



Bias and Fairness in LLMs

What is
Fairness/Bias?



One Word, Four Meanings

You already use “bias” in three other senses. None of them is the one in the headlines.

Statistics

Machine learning

AI in general – Inductive bias

Social science


Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006) for the third sense.



One Word, Four Meanings

You already use “bias” in three other senses. None of them is the one in the headlines.

Statistics



$$\text{bias}(\theta) = E[\theta] - \theta$$

The estimator is systematically off-centre. A property of a procedure against a ground truth.

Machine learning

AI in general – Inductive bias

Social science


Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006) for the third sense.



One Word, Four Meanings

You already use “bias” in three other senses. None of them is the one in the headlines.


Statistics



$$\text{bias}(\theta) = E[\theta] - \theta$$

The estimator is systematically off-centre. A property of a procedure against a ground truth.

Machine learning



$$\text{error} = \text{bias}^2 + \text{variance} + \sigma^2$$

Too simple a model under-fits.
Nothing to do with people.

AI in general – Inductive bias

Social science


Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006) for the third sense.



One Word, Four Meanings

You already use “bias” in three other senses. None of them is the one in the headlines.


Statistics



$$\text{bias}(\theta) = E[\theta] - \theta$$

The estimator is systematically off-centre. A property of a procedure against a ground truth.


Machine learning



$$\text{error} = \text{bias}^2 + \text{variance} + \sigma^2$$

Too simple a model under-fits.
Nothing to do with people.

AI in general – Inductive bias



bias = prior information

“A necessary prerequisite for intelligence.” No prior, no generalisation.

Social science


Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006) for the third sense.



One Word, Four Meanings

You already use “bias” in three other senses. None of them is the one in the headlines.


Statistics



$$\text{bias}(\theta) = E[\theta] - \theta$$

The estimator is systematically off-centre. A property of a procedure against a ground truth.


Machine learning



$$\text{error} = \text{bias}^2 + \text{variance} + \sigma^2$$

Too simple a model under-fits.
Nothing to do with people.

AI in general – Inductive bias



bias = prior information

“A necessary prerequisite for intelligence.” No prior, no generalisation.

Social science

A prior learned from human culture can carry its harms. That is the fourth sense.


Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006) for the third sense.



One Word, Four Meanings

You already use “bias” in three other senses. None of them is the one in the headlines.


Statistics



$$\text{bias}(\theta) = E[\theta] - \theta$$

The estimator is systematically off-centre. A property of a procedure against a ground truth.


Machine learning



$$\text{error} = \text{bias}^2 + \text{variance} + \sigma^2$$

Too simple a model under-fits. Nothing to do with people.


AI in general – Inductive bias



bias = prior information

“A necessary prerequisite for intelligence.” No prior, no generalisation.

Social science



social bias

A systematic skew against a social group. The only one of the four with a victim.

A prior learned from human culture can carry its harms. That is the fourth sense.


Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006) for the third sense.



To Be Intelligent Is to Be Biased

The third sense, taken seriously — because it is the one that makes learning possible.

Which curve?



three observations

Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006). Curves are illustrative.




To Be Intelligent Is to Be Biased

The third sense, taken seriously — because it is the one that makes learning possible.

Which curve?

Infinitely many pass through the same three points.
The data alone does not choose one.



three observations

Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006). Curves are illustrative.



To Be Intelligent Is to Be Biased


The third sense, taken seriously — because it is the one that makes learning possible.

Which curve?

Infinitely many pass through the same three points.
The data alone does not choose one.

The choice is the prior.

“Prefer the smooth one” is an assumption you brought, not a fact you measured.



Caliskan, Bryson & Narayanan, Science (2017), citing Bishop (2006). Curves are illustrative.



To Be Intelligent Is to Be Biased

The third sense, taken seriously — because it is the one that makes learning possible.


Which curve?

Infinitely many pass through the same three points.
The data alone does not choose one.

The choice is the prior.

“Prefer the smooth one” is an assumption you brought, not a fact you measured.

“In AI and machine learning, bias refers to prior information — a necessary prerequisite for intelligence.”



Caliskan, Bryson & Narayanan, *Science* (2017), citing Bishop (2006). Curves are illustrative.



To Be Intelligent Is to Be Biased

The third sense, taken seriously — because it is the one that makes learning possible.


Which curve?

Infinitely many pass through the same three points.
The data alone does not choose one.

The choice is the prior.

“Prefer the smooth one” is an assumption you brought, not a fact you measured.

“In AI and machine learning, bias refers to prior information — a necessary prerequisite for intelligence.”



Caliskan, Bryson & Narayanan, *Science* (2017), citing Bishop (2006). Curves are illustrative.

No prior, no learning. The question is which prior.



To Be Biased Is Not Always to Be Intelligent

Where does the prior come from?

The doctor said that ____ would be late.

The nurse said that ____ would be late.

Caliskan, Bryson & Narayanan, Science (2017). Completions illustrative; the measurement comes later.



To Be Biased Is Not Always to Be Intelligent

Where does the prior come from?

The doctor said that ____ would be late.

he

The nurse said that ____ would be late.

she

Caliskan, Bryson & Narayanan, Science (2017). Completions illustrative; the measurement comes later.



To Be Biased Is Not Always to Be Intelligent

Where does the prior come from?

The doctor said that ____ would be late.

he

The nurse said that ____ would be late.

she

The prior was learned from text people wrote. The precedent it encodes — who does which job — is one we have recognised as harmful.

Caliskan, Bryson & Narayanan, Science (2017). Completions illustrative; the measurement comes later.



To Be Biased Is Not Always to Be Intelligent

Where does the prior come from?

The doctor said that ____ would be late.

he

The nurse said that ____ would be late.

she

The prior was learned from text people wrote. The precedent it encodes — who does which job — is one we have recognised as harmful.

“...bias can be problematic where prior information is derived from precedents known to be harmful.”

Caliskan, Bryson & Narayanan, Science (2017). Completions illustrative; the measurement comes later.



To Be Biased Is Not Always to Be Intelligent

Where does the prior come from?

The doctor said that ____ would be late.

he

The nurse said that ____ would be late.

she

The prior was learned from text people wrote. The precedent it encodes — who does which job — is one we have recognised as harmful.

“...bias can be problematic where prior information is derived from precedents known to be harmful.”

Such a prior is not intelligence. That is the bias we mean today.

Caliskan, Bryson & Narayanan, Science (2017). Completions illustrative; the measurement comes later.



Bias in the Context of Today's Lecture

“an output that is not only **skewed**, but skewed in a way which is **unfair**”

Centre for Data Ethics and Innovation, Review into Bias in Algorithmic Decision-Making (2020)



Bias in the Context of Today's Lecture

“an output that is not only **skewed**, but skewed in a way which is **unfair**”

Centre for Data Ethics and Innovation, Review into Bias in Algorithmic Decision-Making (2020)

skewed

A systematic tilt. A measurable fact about the output — you can count it.



Bias in the Context of Today's Lecture

“an output that is not only **skewed**, but skewed in a way which is **unfair**”

Centre for Data Ethics and Innovation, Review into Bias in Algorithmic Decision-Making (2020)

skewed

A systematic tilt. A measurable fact about the output — you can count it.

unfair

A judgement, not a count. It needs a standard of fairness to be measured against.



Bias in the Context of Today's Lecture

“an output that is not only **skewed**, but skewed in a way which is **unfair**”

Centre for Data Ethics and Innovation, Review into Bias in Algorithmic Decision-Making (2020)

skewed

A systematic tilt. A measurable fact about the output — you can count it.

unfair


A judgement, not a count. It needs a standard of fairness to be measured against.

Skew alone is not bias. Bias presupposes a definition of (un)fairness.



Unfairness, Bias, Discrimination

How the three words nest?



After CDEI (2020), Review into Bias in Algorithmic Decision-Making — redrawn.



Trustworthy Large Language Models



Yatin Nandwani



Unfairness, Bias, Discrimination

How the three words nest?

Synonyms

equity. (FAIRNESS)

impartiality.

justice (FAIRNESS)

fairness *noun* [U] (FAIR TREATMENT)

C1

the quality of treating people equally or in a way that is right or reasonable:

- *He had a real sense of fairness and hated injustice.*
- *The ban on media reporting has made some people question the fairness of the election (= ask whether it was fair).*


(Un)fairness

After CDEI (2020), Review into Bias in Algorithmic Decision-Making — redrawn.



Unfairness, Bias, Discrimination

How the three words nest?



Unfairness

Quality of treating people unequally or in a way that is wrong or unreasonable.

* INEQUALITY

* PARTIALITY

* INJUSTICE

After CDEI (2020), Review into Bias in Algorithmic Decision-Making — redrawn.



Trustworthy Large Language Models




Yatin Nandwani



Unfairness, Bias, Discrimination

How the three words nest?



Unfairness

Quality of treating people unequally or in a way that is wrong or unreasonable.

* INEQUALITY * PARTIALITY * INJUSTICE

Bias


An unfair skew — on any basis: a protected trait, an arbitrary factor, a proxy.

After CDEI (2020), Review into Bias in Algorithmic Decision-Making — redrawn.



Unfairness, Bias, Discrimination

How the three words nest?



Unfairness

No skew required. An opaque process, or no way to appeal a decision, is unfair on its own.

Bias

An unfair skew — on any basis: a protected trait, an arbitrary factor, a proxy.

Discrimination


The legally actionable core: unfair treatment tied to a protected characteristic.

After CDEI (2020), Review into Bias in Algorithmic Decision-Making — redrawn.



Unfairness, Bias, Discrimination

How the three words nest?



Unfairness

No skew required. An opaque process, or no way to appeal a decision, is unfair on its own.

Bias

An unfair skew — on any basis: a protected trait, an arbitrary factor, a proxy.

Discrimination

The legally actionable core: unfair treatment tied to a protected characteristic.

Removing bias does not make a system fair.

After CDEI (2020), Review into Bias in Algorithmic Decision-Making — redrawn.



Older Than Any Algorithm

“Fairness as a concept can be traced to the earliest schools of philosophy.”

Quotes: CDEI / Faculty (2020). Aristotle, *Nicomachean Ethics V*; Rawls, *A Theory of Justice* (1971).



Fairness: Older Than Any Algorithm

“Fairness as a concept can be traced to the earliest schools of philosophy.”

Aristotle

c. 350 BCE

Equals treated equally;
unequals unequally — in
proportion to relevant
difference.

Quotes: CDEI / Faculty (2020). Aristotle, *Nicomachean Ethics* V; Rawls, *A Theory of Justice* (1971).



Fairness: Older Than Any Algorithm

“Fairness as a concept can be traced to the earliest schools of philosophy.”

Aristotle

c. 350 BCE

Equals treated equally;
unequals unequally — in
proportion to relevant
difference.

Rawls

1971

“Justice as fairness”: rules you
would accept without knowing
your place in society.

Quotes: CDEI / Faculty (2020). Aristotle, *Nicomachean Ethics V*; Rawls, *A Theory of Justice* (1971).



Trustworthy Large Language Models



Yatin Nandwani

Fairness: Older Than Any Algorithm

“Fairness as a concept can be traced to the earliest schools of philosophy.”

Aristotle

c. 350 BCE

Equals treated equally;
unequals unequally — in
proportion to relevant
difference.

Rawls

1971

“Justice as fairness”: rules you
would accept without knowing
your place in society.

Politics & policy

today

“Along with justice, perhaps
the most debated issue within
politics and social policy.”

Quotes: CDEI / Faculty (2020). Aristotle, Nicomachean Ethics V; Rawls, A Theory of Justice (1971).



Trustworthy Large Language Models



Yatin Nandwani

Fairness: Older Than Any Algorithm

“Fairness as a concept can be traced to the earliest schools of philosophy.”

Aristotle

c. 350 BCE

Equals treated equally;
unequals unequally — in
proportion to relevant
difference.

Rawls

1971

“Justice as fairness”: rules you
would accept without knowing
your place in society.

Politics & policy

today

“Along with justice, perhaps
the most debated issue within
politics and social policy.”

Historically an abstract quality, not a measurable one.

Quotes: CDEI / Faculty (2020). Aristotle, *Nicomachean Ethics V*; Rawls, *A Theory of Justice* (1971).



Trustworthy Large Language Models



Yatin Nandwani


Is Fairness a Law of Nature?

**Is fairness a law of nature —
or something we made up?**



Is Fairness a Law of Nature?

Two capuchin monkeys, one experimenter, two kinds of reward.



Both do the same task.

Brosnan & de Waal, "Monkeys reject unequal pay", Nature 425 (2003).




Trustworthy Large Language Models




Yatin Nandwani

Is Fairness a Law of Nature?

Two capuchin monkeys, one experimenter, two kinds of reward.



Both do the same task.




Both are paid a cucumber slice. Both content.

Brosnan & de Waal, "Monkeys reject unequal pay", Nature 425 (2003).




Is Fairness a Law of Nature?


Two capuchin monkeys, one experimenter, two kinds of reward.



Both do the same task.



Both are paid a cucumber slice. Both content.




Same task — one is now paid a grape.

Brosnan & de Waal, "Monkeys reject unequal pay", Nature 425 (2003).




Is Fairness a Law of Nature?


Two capuchin monkeys, one experimenter, two kinds of reward.




Both do the same task.



Both are paid a cucumber slice. Both content.



Same task — one is now paid a grape.




The other refuses the cucumber.

Brosnan & de Waal, "Monkeys reject unequal pay", Nature 425 (2003).




Is Fairness a Law of Nature?


Two capuchin monkeys, one experimenter, two kinds of reward.




Both do the same task.



Both are paid a cucumber slice. Both content.



Same task — one is now paid a grape.



The other refuses the cucumber.

Nature has the impulse.


A sense of unequal pay — or frustration at a visible grape? Contested since 2003.

Brosnan & de Waal, "Monkeys reject unequal pay", Nature 425 (2003).




Is Fairness a Law of Nature?


Two capuchin monkeys, one experimenter, two kinds of reward.




Both do the same task.



Both are paid a cucumber slice. Both content.



Same task — one is now paid a grape.



The other refuses the cucumber.

Nature has the impulse.

A sense of unequal pay — or frustration at a visible grape? Contested since 2003.

Nature has no definition.

No species negotiates what counts as “the same work”. That argument is ours.

Brosnan & de Waal, “Monkeys reject unequal pay”, Nature 425 (2003).




Trustworthy Large Language Models




Yatin Nandwani

Is Fairness a Law of Nature?


Two capuchin monkeys, one experimenter, two kinds of reward.




Both do the same task.



Both are paid a cucumber slice. Both content.



Same task — one is now paid a grape.



The other refuses the cucumber.

Nature has the impulse.

A sense of unequal pay — or frustration at a visible grape? Contested since 2003.

Nature has no definition.

No species negotiates what counts as “the same work”. That argument is ours.

Fairness is a man-made construct — and our notion of it keeps evolving.

Brosnan & de Waal, “Monkeys reject unequal pay”, Nature 425 (2003).



The Standard of Fairness Moves

Every one of these was, at the time, simply how a bank or a ballot worked.

1928

United Kingdom

Women vote on the same terms as men. Ten years earlier the age bar was 30 — and property-qualified.

Representation of the People Act 1928; Equal Credit Opportunity Act 1974; Equality Act 2010.



The Standard of Fairness Moves

Every one of these was, at the time, simply how a bank or a ballot worked.

1928

United Kingdom

Women vote on the same terms as men. Ten years earlier the age bar was 30 — and property-qualified.

1974

United States

The Equal Credit Opportunity Act. Until then a bank could require a woman to bring a male co-signer.

Representation of the People Act 1928; Equal Credit Opportunity Act 1974; Equality Act 2010.



The Standard of Fairness Moves

Every one of these was, at the time, simply how a bank or a ballot worked.

1928

United Kingdom

Women vote on the same terms as men. Ten years earlier the age bar was 30 — and property-qualified.

1974

United States

The Equal Credit Opportunity Act. Until then a bank could require a woman to bring a male co-signer.

2010

United Kingdom

The Equality Act names nine protected characteristics and makes indirect discrimination unlawful.

Representation of the People Act 1928; Equal Credit Opportunity Act 1974; Equality Act 2010.



The Standard of Fairness Moves

Every one of these was, at the time, simply how a bank or a ballot worked.

1928

United Kingdom

Women vote on the same terms as men. Ten years earlier the age bar was 30 — and property-qualified.

1974

United States

The Equal Credit Opportunity Act. Until then a bank could require a woman to bring a male co-signer.

2010

United Kingdom


The Equality Act names nine protected characteristics and makes indirect discrimination unlawful.

Representation of the People Act 1928; Equal Credit Opportunity Act 1974; Equality Act 2010.

What is considered “fair” keeps evolving.



Systems to Ensure Fairness



Societies have never trusted a decision-maker to be fair. They build procedure instead.


What we built

- A written rule, fixed before the case
- A stated reason for the decision
- A right to be heard, and to appeal
- Someone independent who can audit it

Procedural safeguards as characterised in CDEI (2020).



Systems to Ensure Fairness



Societies have never trusted a decision-maker to be fair. They build procedure instead.

What we built

- A written rule, fixed before the case
- A stated reason for the decision
- A right to be heard, and to appeal
- Someone independent who can audit it


Why it still leaks

- Discretion hides in the gaps between rules
- The stated reason need not be the real one
- An appeal needs someone to notice first
- The auditor sees outcomes, not intentions

Procedural safeguards as characterised in CDEI (2020).



Systems to Ensure Fairness



Societies have never trusted a decision-maker to be fair. They build procedure instead.

What we built

- A written rule, fixed before the case
- A stated reason for the decision
- A right to be heard, and to appeal
- Someone independent who can audit it

Why it still leaks


- Discretion hides in the gaps between rules
- The stated reason need not be the real one
- An appeal needs someone to notice first
- The auditor sees outcomes, not intentions

Never foolproof — but someone decided, could be asked why, and could be overruled.

Procedural safeguards as characterised in CDEI (2020).



The Decision Moved to an Algorithm




COMPAS

2016

Risk scores used in US pretrial decisions. Among defendants who did not go on to reoffend, ~45% of Black and ~23% of white defendants had been labelled high-risk.



The Decision Moved to an Algorithm



COMPAS

2016

Risk scores used in US pretrial decisions. Among defendants who did not go on to reoffend, ~45% of Black and ~23% of white defendants had been labelled high-risk.


Amazon's CV screener

2014–17

Trained on ten years of the company's own hires. It learned to penalise the word “women’s” — as in women’s chess club captain. Scrapped.



The Decision Moved to an Algorithm



COMPAS

2016

Risk scores used in US pretrial decisions. Among defendants who did not go on to reoffend, ~45% of Black and ~23% of white defendants had been labelled high-risk.

Amazon's CV screener

2014–17

Trained on ten years of the company's own hires. It learned to penalise the word “women’s” — as in women’s chess club captain. Scrapped.


Workday

in court now

One screening service, many employers. The lawsuit in the headline you saw ten minutes ago.



The Decision Moved to an Algorithm



COMPAS

2016

Risk scores used in US pretrial decisions. Among defendants who did not go on to reoffend, ~45% of Black and ~23% of white defendants had been labelled high-risk.

Amazon's CV screener

2014–17

Trained on ten years of the company's own hires. It learned to penalise the word “women's” — as in women's chess club captain. Scrapped.

Workday

in court now

One screening service, many employers. The lawsuit in the headline you saw ten minutes ago.


ProPublica (2016), figures via Chouldechova (2017); Reuters (2018), Amazon; Reuters (2026), Workday.

Not one of these is a language model. The problem is older than the technology.



One Person, or One Algorithm

What changes when the same skew moves from a desk into a deployment?




One desk. Hundreds of decisions.




One Person, or One Algorithm

What changes when the same skew moves from a desk into a deployment?



One desk. Hundreds of decisions.




Every deployment. Millions of decisions.




One Person, or One Algorithm

What changes when the same skew moves from a desk into a deployment?



One desk. Hundreds of decisions.



Every deployment. Millions of decisions.


Reach

The same skew, applied everywhere the algorithm is used, at once.




One Person, or One Algorithm

What changes when the same skew moves from a desk into a deployment?



One desk. Hundreds of decisions.



Every deployment. Millions of decisions.

Reach

The same skew, applied everywhere the algorithm is used, at once.

**Extensive reach calls for a system to ensure algorithms are fair —
Algorithmic Fairness.**



Algorithmic Fairness: A Framework

The first question to ask of any decision-making algorithm.



Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision-Making.



Algorithmic Fairness: A Framework

The first question to ask of any decision-making algorithm.




Was the person treated fairly in the process by which the decision was made?

Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision-Making.



Procedural Fairness

Fair treatment of the person within the process — not only inside the model.




CDEI (2020): procedural fairness “encompasses wider elements of the decision-making process” than the algorithm’s design.



Procedural Fairness

Fair treatment of the person within the process — not only inside the model.




- Can the person see how the decision was made?

CDEI (2020): procedural fairness “encompasses wider elements of the decision-making process” than the algorithm’s design.



Procedural Fairness

Fair treatment of the person within the process — not only inside the model.




- Can the person see how the decision was made?
- Can they question or rebut it?

CDEI (2020): procedural fairness “encompasses wider elements of the decision-making process” than the algorithm’s design.



Procedural Fairness

Fair treatment of the person within the process — not only inside the model.




- Can the person see how the decision was made?
- Can they question or rebut it?
- Can they seek redress if it was wrong?

CDEI (2020): procedural fairness “encompasses wider elements of the decision-making process” than the algorithm’s design.



Procedural Fairness

Fair treatment of the person within the process — not only inside the model.



- Can the person see how the decision was made?
- Can they question or rebut it?
- Can they seek redress if it was wrong?

The whole decision cycle in the real world — before, during and after the algorithm runs.

CDEI (2020): procedural fairness “encompasses wider elements of the decision-making process” than the algorithm’s design.



Fairness Through Unawareness

The first thing every engineer proposes: do not give the model the sensitive attributes.

income


employment

existing debt

postcode

~~gender~~

~~race~~



CDEI (2020). Postcode–race is the review’s own example; Amazon’s screener never saw gender and learned “women’s”.



Fairness Through Unawareness

The first thing every engineer proposes: do not give the model the sensitive attributes.

income


employment

existing debt

postcode

gender

race



The premise

If the model never sees gender or race, it cannot discriminate on them.

CDEI (2020). Postcode–race is the review’s own example; Amazon’s screener never saw gender and learned “women’s”.



Fairness Through Unawareness

The first thing every engineer proposes: do not give the model the sensitive attributes.

income

employment


existing debt

postcode

~~gender~~

~~race~~

proxy



The premise

If the model never sees gender or race, it cannot discriminate on them.

CDEI (2020). Postcode–race is the review’s own example; Amazon’s screener never saw gender and learned “women’s”.



Fairness Through Unawareness

The first thing every engineer proposes: do not give the model the sensitive attributes.

income

employment


existing debt

postcode

gender

race

proxy



The premise

If the model never sees gender or race, it cannot discriminate on them.

The gap

Postcode can predict race. The model uses postcode. Nothing was checked at the output.

CDEI (2020). Postcode–race is the review’s own example; Amazon’s screener never saw gender and learned “women’s”.



Fairness Through Unawareness

The first thing every engineer proposes: do not give the model the sensitive attributes.

income

employment


existing debt

postcode

gender

race

proxy



The premise

If the model never sees gender or race, it cannot discriminate on them.

The gap

Postcode can predict race. The model uses postcode. Nothing was checked at the output.

Procedural: it constrains the inputs and says nothing about the outcome.

CDEI (2020). Postcode–race is the review’s own example; Amazon’s screener never saw gender and learned “women’s”.



Algorithmic Fairness: A Framework

The second question: not the process, the result.




Was the person treated fairly in the process by which the decision was made?

Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision-Making.



Algorithmic Fairness: A Framework

The second question: not the process, the result.




Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision-Making.



Algorithmic Fairness: A Framework

The second question: not the process, the result.




Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision-Making.



Algorithmic Fairness: A Framework

The second question: not the process, the result.




Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision-Making.



Algorithmic Fairness: A Framework

The second question: not the process, the result.




Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision-Making.



Outcome Fairness: Individual

Similar individuals should receive similar decisions.




Dwork, Hardt, Pitassi, Reingold & Zemel, "Fairness Through Awareness" (ITCS 2012).



Outcome Fairness: Individual

Similar individuals should receive similar decisions.




Similar — by what measure?

Dwork, Hardt, Pitassi, Reingold & Zemel, “Fairness Through Awareness” (ITCS 2012).



Outcome Fairness: Individual

Similar individuals should receive similar decisions.



Similar — by what measure?

Choosing the similarity measure is itself a judgement about which differences should matter. It reintroduces the problem it was meant to avoid.

Dwork, Hardt, Pitassi, Reingold & Zemel, "Fairness Through Awareness" (ITCS 2012).




Outcome Fairness: Group

Different groups should be treated alike. Which groups?



G_1



G_2

Equality Act 2010, s.4. India, Art. 15: religion, race, caste, sex, place of birth.




Outcome Fairness: Group

Different groups should be treated alike. Which groups?



G_1



G_2

Any partition of the population — or the groups society has historically treated unfairly?

Equality Act 2010, s.4. India, Art. 15: religion, race, caste, sex, place of birth.




Outcome Fairness: Group

Different groups should be treated alike. Which groups?



G_1



G_2

Any partition of the population — or the groups society has historically treated unfairly?

Protected characteristics — UK Equality Act 2010

age

gender reassignment

pregnancy and maternity

religion or belief

sexual orientation

disability

marriage or civil partnership

race

sex

Equality Act 2010, s.4. India, Art. 15: religion, race, caste, sex, place of birth.




Outcome Fairness: Group

Different groups should be treated alike. Which groups?



G_1



G_2

Any partition of the population — or the groups society has historically treated unfairly?

For groups G_1 and G_2 , algorithm A should decide “similarly”. How do we operationalise that?

Equality Act 2010, s.4. India, Art. 15: religion, race, caste, sex, place of birth.

Protected characteristics — UK Equality Act 2010

age

gender reassignment

pregnancy and maternity

religion or belief

sexual orientation

disability

marriage or civil partnership

race

sex



Group Fairness: Demographic Parity

The first way to operationalise “decides similarly”.

$$P(\text{decision} = \text{yes} \mid G_1) = P(\text{decision} = \text{yes} \mid G_2)$$

Background

the same acceptance rate — or the same error rate — in both groups

Also called independence / statistical parity. Buolamwini & Gebru (2018) on face analysis.



Group Fairness: Demographic Parity

The first way to operationalise “decides similarly”.

$$P(\text{decision} = \text{yes} \mid G_1) = P(\text{decision} = \text{yes} \mid G_2)$$

the same acceptance rate — or the same error rate — in both groups

Where it fits

Face detection. The error rate should not depend on skin tone — there is no legitimate reason for it to.

Also called independence / statistical parity. Buolamwini & Gebru (2018) on face analysis.



Group Fairness: Demographic Parity

The first way to operationalise “decides similarly”.

$$P(\text{decision} = \text{yes} \mid G_1) = P(\text{decision} = \text{yes} \mid G_2)$$

the same acceptance rate — or the same error rate — in both groups

Where it fits

Face detection. The error rate should not depend on skin tone — there is no legitimate reason for it to.

Where it does not

Loan approval. Gender should play no role — but if repayment rates differ between groups, even a perfect predictor fails this test.

Also called independence / statistical parity. Buolamwini & Gebru (2018) on face analysis.



Group Fairness: Demographic Parity

The first way to operationalise “decides similarly”.

$$P(\text{decision} = \text{yes} \mid G_1) = P(\text{decision} = \text{yes} \mid G_2)$$

the same acceptance rate — or the same error rate — in both groups

Where it fits

Face detection. The error rate should not depend on skin tone — there is no legitimate reason for it to.

Where it does not

Loan approval. Gender should play no role — but if repayment rates differ between groups, even a perfect predictor fails this test.



So the definition must take into account what actually happened.

Also called independence / statistical parity. Buolamwini & Gebru (2018) on face analysis.



Group Fairness: Equalised Odds

Loan approval, with hindsight: who actually repaid, and who was approved.



Hardt, Price & Srebro, NeurIPS 2016. Also called separation.



Group Fairness: Equalised Odds

Loan approval, with hindsight: who actually repaid, and who was approved.

		<i>model's decision</i>	
		approve	refuse
<i>what actually happened</i>	repaid	TP	FN
	defaulted	FP	TN

Hardt, Price & Srebro, NeurIPS 2016. Also called separation.



Group Fairness: Equalised Odds

Loan approval, with hindsight: who actually repaid, and who was approved.

		<i>model's decision</i>	
		approve	refuse
<i>what actually happened</i>	repaid	TP	FN
	defaulted	FP	TN

True-positive rate

$$\text{TPR} = \text{TP} / (\text{TP} + \text{FN})$$

of those who would repay, the fraction approved

False-positive rate

$$\text{FPR} = \text{FP} / (\text{FP} + \text{TN})$$

of those who would default, the fraction approved

Hardt, Price & Srebro, NeurIPS 2016. Also called separation.



Group Fairness: Equalised Odds

Loan approval, with hindsight: who actually repaid, and who was approved.

		<i>model's decision</i>	
		approve	refuse
<i>what actually happened</i>	repaid	TP	FN
	defaulted	FP	TN

True-positive rate

$$\text{TPR} = \text{TP} / (\text{TP} + \text{FN})$$

of those who would repay, the fraction approved

False-positive rate

$$\text{FPR} = \text{FP} / (\text{FP} + \text{TN})$$

of those who would default, the fraction approved

$$\text{TPR}(G_1) = \text{TPR}(G_2) \quad \text{and} \quad \text{FPR}(G_1) = \text{FPR}(G_2)$$

Hardt, Price & Srebro, NeurIPS 2016. Also called separation.



Group Fairness: Equalised Odds

Loan approval, with hindsight: who actually repaid, and who was approved.

		<i>model's decision</i>	
		approve	refuse
<i>what actually happened</i>	repaid	TP	FN
	defaulted	FP	TN

True-positive rate

$$\text{TPR} = \text{TP} / (\text{TP} + \text{FN})$$

of those who would repay, the fraction approved

False-positive rate

$$\text{FPR} = \text{FP} / (\text{FP} + \text{TN})$$

of those who would default, the fraction approved


$$\text{TPR}(G_1) = \text{TPR}(G_2) \quad \text{and} \quad \text{FPR}(G_1) = \text{FPR}(G_2)$$

Equal chances of a right — and a wrong — decision in every group.

Hardt, Price & Srebro, NeurIPS 2016. Also called separation.




Algorithmic Fairness: A Framework



Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision Making




Algorithmic Fairness: In the Era of LLMs



Organising frame after CDEI (2020), Review into Bias in Algorithmic Decision Making



Algorithmic Fairness: In the Era of LLMs




Before

One model, one decision, one audit.




Algorithmic Fairness: In the Era of LLMs

HELM



Before

One model, one decision, one audit.




Now

One model, every decision. One skew, everywhere.




Algorithmic Fairness: In the Era of LLMs



Before

One model, one decision, one audit.

- One model, every decision — a single skew propagates into every downstream use at once




Now

One model, every decision. One skew, everywhere.




Algorithmic Fairness: In the Era of LLMs



Before

One model, one decision, one audit.

- One model, every decision — a single skew propagates into every downstream use at once
- The model is now the judge — bias in the evaluator is invisible to every metric built on its scores




Now

One model, every decision. One skew, everywhere.



Algorithmic Fairness: In the Era of LLMs




Before

One model, one decision, one audit.

- One model, every decision — a single skew propagates into every downstream use at once
- The model is now the judge — bias in the evaluator is invisible to every metric built on its scores
- Nobody sees the decision boundary — no coefficients to audit, just weights and a prompt

The decision layer is now a language model.




Now

One model, every decision. One skew, everywhere.



Not Only Decisions

Share of English web pages showing signs of AI editing or authorship.




Pew Research Center (2026): 490,000 English pages from Common Crawl, Jan 2021 – Jul 2026, scored with Pangram.
Curves are the published chart; axes and labels redrawn.



Not Only Decisions

Share of English web pages showing signs of AI editing or authorship.



■ 35%


of pages published after ChatGPT show signs of AI authorship

Pew Research Center (2026): 490,000 English pages from Common Crawl, Jan 2021 – Jul 2026, scored with Pangram.
Curves are the published chart; axes and labels redrawn.



Not Only Decisions

Share of English web pages showing signs of AI editing or authorship.



■ **35%**

of pages published after ChatGPT show signs of AI authorship

■ **9.6%**


of all pages sampled — old and new together

Pew Research Center (2026): 490,000 English pages from Common Crawl, Jan 2021 – Jul 2026, scored with Pangram.
Curves are the published chart; axes and labels redrawn.



Not Only Decisions

Share of English web pages showing signs of AI editing or authorship.



■ **35%**

of pages published after ChatGPT show signs of AI authorship

■ **9.6%**

of all pages sampled — old and new together

It is text that people read.

The output is no longer only a decision.

Pew Research Center (2026): 490,000 English pages from Common Crawl, Jan 2021 – Jul 2026, scored with Pangram.
Curves are the published chart; axes and labels redrawn.



Two Families of Harm

What kind of damage does the skew do? The two halves need different tools.

REPRESENTATIONAL

How groups are portrayed.

ALLOCATIONAL

How resources and opportunities are distributed.

Gallegos et al. (2024), Table 1 — after Crawford, “The Trouble with Bias,” NeurIPS 2017 keynote



Two Families of Harm

What kind of damage does the skew do? The two halves need different tools.

REPRESENTATIONAL

How groups are portrayed.

Derogatory language — slurs, pejoratives

Disparate performance — works worse for some

Erasure — a group absent from the output

Exclusionary norms — one group as default

Misrepresentation — a distorted picture

Stereotyping — fixed attributes for a group

Toxicity — hostile content aimed at a group

ALLOCATIONAL

How resources and opportunities are distributed.

Gallegos et al. (2024), Table 1 — after Crawford, “The Trouble with Bias,” NeurIPS 2017 keynote



Two Families of Harm

What kind of damage does the skew do? The two halves need different tools.

REPRESENTATIONAL

How groups are portrayed.

Derogatory language — slurs, pejoratives

Disparate performance — works worse for some

Erasure — a group absent from the output

Exclusionary norms — one group as default

Misrepresentation — a distorted picture

Stereotyping — fixed attributes for a group

Toxicity — hostile content aimed at a group

ALLOCATIONAL

How resources and opportunities are distributed.

Direct discrimination — worse treatment because of a protected characteristic

Indirect discrimination — a neutral rule that disadvantages a group, via proxies

Two categories against seven — and yet almost the entire fairness-metrics literature you just met was built for this column.

Gallegos et al. (2024), Table 1 — after Crawford, “The Trouble with Bias,” NeurIPS 2017 keynote



Why That Split Is the Hinge

What every criterion so far required — and what a sentence does not have.



Why That Split Is the Hinge

What every criterion so far required — and what a sentence does not have.

Allocational harm has a confusion matrix.

- It needs a decision, a group and an outcome
- Given those three, parity, odds and calibration are computable



Why That Split Is the Hinge

What every criterion so far required — and what a sentence does not have.

Allocational harm has a confusion matrix.

- It needs a decision, a group and an outcome
- Given those three, parity, odds and calibration are computable

Representational harm does not.

- No decision to audit, no positive class, no base rate
- A sentence is demeaning or it is not



Why That Split Is the Hinge

What every criterion so far required — and what a sentence does not have.

Allocational harm has a confusion matrix.

- It needs a decision, a group and an outcome
- Given those three, parity, odds and calibration are computable

Representational harm does not.

- No decision to audit, no positive class, no base rate
- A sentence is demeaning or it is not

So the metrics do not transfer.

- Equalised odds cannot be computed on a paragraph
- Measuring bias in text needs a different apparatus



Why That Split Is the Hinge

What every criterion so far required — and what a sentence does not have.

Allocational harm has a confusion matrix.

- It needs a decision, a group and an outcome
- Given those three, parity, odds and calibration are computable

Representational harm does not.

- No decision to audit, no positive class, no base rate
- A sentence is demeaning or it is not

So the metrics do not transfer.

- Equalised odds cannot be computed on a paragraph
- Measuring bias in text needs a different apparatus

That apparatus is Lecture 2.



Sources of Bias in LLMs

1


Training corpus →

1 • Training corpus

The web is not a sample of humanity: over- and under-representation, and filtering nobody documented.



Sources of Bias in LLMs



1 • Training corpus

The web is not a sample of humanity: over- and under-representation, and filtering nobody documented.

2 • Training objective

The loss rewards fitting co-occurrence — and amplifies it. Nothing in it opposes bias.



Sources of Bias in LLMs



1 • Training corpus

The web is not a sample of humanity: over- and under-representation, and filtering nobody documented.

2 • Training objective


The loss rewards fitting co-occurrence — and amplifies it. Nothing in it opposes bias.

3 • Alignment — SFT / RLHF

A narrow annotator pool. Optimising to an average preference can collapse minority views.



Sources of Bias in LLMs



1 • Training corpus

The web is not a sample of humanity: over- and under-representation, and filtering nobody documented.

2 • Training objective

The loss rewards fitting co-occurrence — and amplifies it. Nothing in it opposes bias.

3 • Alignment — SFT / RLHF

A narrow annotator pool. Optimising to an average preference can collapse minority views.

4 • Prompt & deployment

System prompt, persona, in-context examples, option order. None of it is in the weights.



Sources of Bias in LLMs



1 • Training corpus

The web is not a sample of humanity: over- and under-representation, and filtering nobody documented.

2 • Training objective

The loss rewards fitting co-occurrence — and amplifies it. Nothing in it opposes bias.

3 • Alignment — SFT / RLHF

A narrow annotator pool. Optimising to an average preference can collapse minority views.

4 • Prompt & deployment

System prompt, persona, in-context examples, option order. None of it is in the weights.

5 • Evaluation & feedback

Benchmarks encode their authors' assumptions — and output re-enters tomorrow's corpus.

Five entry points — five different fixes.



Sources of Bias in LLMs



1 · Training corpus

The web is not a sample of humanity: over- and under-representation, and filtering nobody documented.



Size Doesn't Guarantee Diversity

Common Crawl is “petabytes of data collected over 8 years.” Surely that is everyone?

Four filters narrow the population — all of them before you write a line of training code.

1 Who has internet access Over-represents younger users and developed countries.



Size Doesn't Guarantee Diversity

Common Crawl is “petabytes of data collected over 8 years.” Surely that is everyone?

Four filters narrow the population — all of them before you write a line of training code.

- 1 Who has internet access** Over-represents younger users and developed countries.
- 2 Which subsample is scraped** GPT-2's data came from links posted to Reddit. 67% of US Reddit users are men; 64% are 18–29. Only 8.8–15% of Wikipedians are women or girls.



Size Doesn't Guarantee Diversity

Common Crawl is “petabytes of data collected over 8 years.” Surely that is everyone?

Four filters narrow the population — all of them before you write a line of training code.

- 1 Who has internet access** Over-represents younger users and developed countries.
- 2 Which subsample is scraped** GPT-2's data came from links posted to Reddit. 67% of US Reddit users are men; 64% are 18–29. Only 8.8–15% of Wikipedians are women or girls.
- 3 Who is allowed to stay** People receiving death threats had their accounts suspended while the threatening accounts persisted.



Size Doesn't Guarantee Diversity

Common Crawl is “petabytes of data collected over 8 years.” Surely that is everyone?

Four filters narrow the population — all of them before you write a line of training code.

- 1 Who has internet access** Over-represents younger users and developed countries.
- 2 Which subsample is scraped** GPT-2's data came from links posted to Reddit. 67% of US Reddit users are men; 64% are 18–29. Only 8.8–15% of Wikipedians are women or girls.
- 3 Who is allowed to stay** People receiving death threats had their accounts suspended while the threatening accounts persisted.
- 4 Which alternative spaces get crawled** Older adults' anti-ageist blogs have few inbound links — so they are less likely to be found at all.



Size Doesn't Guarantee Diversity

Common Crawl is “petabytes of data collected over 8 years.” Surely that is everyone?

Four filters narrow the population — all of them before you write a line of training code.

- 1 Who has internet access** Over-represents younger users and developed countries.
- 2 Which subsample is scraped** GPT-2's data came from links posted to Reddit. 67% of US Reddit users are men; 64% are 18–29. Only 8.8–15% of Wikipedians are women or girls.
- 3 Who is allowed to stay** People receiving death threats had their accounts suspended while the threatening accounts persisted.
- 4 Which alternative spaces get crawled** Older adults' anti-ageist blogs have few inbound links — so they are less likely to be found at all.

A larger sample of a narrowed population is still a sample of a narrowed population.



Built for English, Used on Everyone

The population whose text trains the model is not the population it is used on.

Whose text trains it

- English, and a few other high-resource languages
- “Similar large-scale models aren’t being produced for Dhivehi or Sudanese Arabic”
- Written by whoever was online, posting, and allowed to stay

Bender, Gebru, McMillan-Major & Shmitchell (2021), “On the Dangers of Stochastic Parrots”, §2 and §4.1.



Built for English, Used on Everyone

The population whose text trains the model is not the population it is used on.

Whose text trains it

- English, and a few other high-resource languages
- “Similar large-scale models aren’t being produced for Dhivehi or Sudanese Arabic”
- Written by whoever was online, posting, and allowed to stay

Who it is used on

- Speakers of languages with no comparable model
- Reached through translation, search and assistants built on it
- People who supplied none of the training text

Bender, Gebru, McMillan-Major & Shmitchell (2021), “On the Dangers of Stochastic Parrots”, §2 and §4.1.



Built for English, Used on Everyone

The population whose text trains the model is not the population it is used on.

Whose text trains it

- English, and a few other high-resource languages
- “Similar large-scale models aren’t being produced for Dhivehi or Sudanese Arabic”
- Written by whoever was online, posting, and allowed to stay

Who it is used on

- Speakers of languages with no comparable model
- Reached through translation, search and assistants built on it
- People who supplied none of the training text

Selection bias: training population $\neq$ deployment population.

Bender, Gebru, McMillan-Major & Shmitchell (2021), “On the Dangers of Stochastic Parrots”, §2 and §4.1.



And Then We “Clean” It

The last narrowing is the one we perform deliberately, believing it to be hygiene.

GPT-3

Common Crawl, kept only where a classifier judged a page to resemble those linked from Reddit.

Bender et al. (2021), §4.



And Then We “Clean” It

The last narrowing is the one we perform deliberately, believing it to be hygiene.

GPT-3

Common Crawl, kept only where a classifier judged a page to resemble those linked from Reddit.

C4

Any page containing one of ~400 “Dirty, Naughty, Obscene or Otherwise Bad Words” was discarded.

Bender et al. (2021), §4.



And Then We “Clean” It

The last narrowing is the one we perform deliberately, believing it to be hygiene.

GPT-3

Common Crawl, kept only where a classifier judged a page to resemble those linked from Reddit.

C4

Any page containing one of ~400 “Dirty, Naughty, Obscene or Otherwise Bad Words” was discarded.

“What is unmeasured (and thus unknown) is what else it filtered out.”

Bender et al. (2021), §4.



And Then We “Clean” It

The last narrowing is the one we perform deliberately, believing it to be hygiene.

GPT-3

Common Crawl, kept only where a classifier judged a page to resemble those linked from Reddit.

C4

Any page containing one of ~400 “Dirty, Naughty, Obscene or Otherwise Bad Words” was discarded.

“What is unmeasured (and thus unknown) is what else it filtered out.”

The word list includes “twink” — so the filter removes spaces built by and for LGBTQ people.

Bender et al. (2021), §4.



And Then We “Clean” It

The last narrowing is the one we perform deliberately, believing it to be hygiene.

GPT-3

Common Crawl, kept only where a classifier judged a page to resemble those linked from Reddit.

C4

Any page containing one of ~400 “Dirty, Naughty, Obscene or Otherwise Bad Words” was discarded.

“What is unmeasured (and thus unknown) is what else it filtered out.”

The word list includes “twink” — so the filter removes spaces built by and for LGBTQ people.

Cleaning is a value judgement, applied at scale and undocumented.

Bender et al. (2021), §4.



How Would We Show the Bias Is in the Text?

We cannot interview a corpus. But the corpus left something behind we can measure.


The text

hundreds of billions of tokens




How Would We Show the Bias Is in the Text?

We cannot interview a corpus. But the corpus left something behind we can measure.



How Would We Show the Bias Is in the Text?

We cannot interview a corpus. But the corpus left something behind we can measure.




Whatever regularity the text carried – is it visible in the embeddings?



How Would We Show the Bias Is in the Text?

We cannot interview a corpus. But the corpus left something behind we can measure.



Whatever regularity the text carried – is it visible in the embeddings?

So: probe the embeddings. With what instrument?



How Do We Test Implicit Bias in People?


The same problem: you cannot get at it by asking.

Screenshots: implicit.harvard.edu. The test: Greenwald, McGhee & Schwartz (1998).



How Do We Test Implicit Bias in People?

The same problem: you cannot get at it by asking.



Project Implicit

Harvard's public site has run the Implicit Association Test online since 1998, across a menu of topics — race, gender and science, age, disability.

Screenshots: implicit.harvard.edu. The test: Greenwald, McGhee & Schwartz (1998).





Trustworthy Large Language Models



Yatin Nandwani

How Do We Test Implicit Bias in People?

The same problem: you cannot get at it by asking.



Screenshots: implicit.harvard.edu. The test: Greenwald, McGhee & Schwartz (1998).

Project Implicit

Harvard's public site has run the **Implicit Association Test** online since 1998, across a menu of topics — race, gender and science, age, disability.



Why an indirect test

An association a person will not report — or does not know they hold — cannot be reached by a questionnaire. So measure something they cannot control: how fast they respond.



How Do We Test Implicit Bias in People?

The same problem: you cannot get at it by asking.



Screenshots: implicit.harvard.edu. The test: Greenwald, McGhee & Schwartz (1998).

Project Implicit

Harvard's public site has run the Implicit Association Test online since 1998, across a menu of topics — race, gender and science, age, disability.

Why an indirect test

An association a person will not report — or does not know they hold — cannot be reached by a questionnaire. So measure something they cannot control: how fast they respond.

Reaction time is the measure.



The Implicit Association Test

Sort words into groups using two keys, as fast as you can.

Target groups

Male

Ben · Paul · Daniel · John · Jeffrey

Female

Rebecca · Michelle · Emily · Julia · Anna

Greenwald, McGhee & Schwartz (1998).



Trustworthy Large Language Models



Yatin Nandwani



The Implicit Association Test

Sort words into groups using two keys, as fast as you can.

Target groups

Male

Ben · Paul · Daniel · John · Jeffrey

Female

Rebecca · Michelle · Emily · Julia · Anna

Attributes

Career

Career · Corporation · Salary · Office · Professional · Business

Family

Wedding · Marriage · Parents · Relatives · Home · Children

Greenwald, McGhee & Schwartz (1998).



Trustworthy Large Language Models



Yatin Nandwani

The Implicit Association Test

Sort words into groups using two keys, as fast as you can.

Target groups

Male

Ben · Paul · Daniel · John · Jeffrey

Female

Rebecca · Michelle · Emily · Julia · Anna

Attributes

Career

Career · Corporation · Salary · Office · Professional · Business

Family

Wedding · Marriage · Parents · Relatives · Home · Children

Two target groups, two attributes, two keys — and five blocks of trials.

Greenwald, McGhee & Schwartz (1998).



Trustworthy Large Language Models



Yatin Nandwani

The Implicit Association Test

Sort words into groups using two keys, as fast as you can.

Target groups

Male

Ben · Paul · Daniel · John · Jeffrey

Female

Rebecca · Michelle · Emily · Julia · Anna

Attributes

Career

Career · Corporation · Salary · Office · Professional · Business

Family

Wedding · Marriage · Parents · Relatives · Home · Children

Two target groups, two attributes, two keys — and five blocks of trials.

The measurement: how long you take when the pairing matches an association.

Greenwald, McGhee & Schwartz (1998).



The Implicit Association Test

	Block 1
Task description	Initial attribute discrimination
Example trial	Family Career Wedding
Key	 E  I

Block 1

E for family words, I for career words. Pure practice — you are just reading nouns.



The Implicit Association Test

	Block 1	Block 2
Task description	Initial attribute discrimination	Associated target discrimination
Example trial	Family Career <u>Family</u> Wedding	Male Female <u>Male</u> Daniel
Key	  E I	  <u>E</u> I

Block 2

E for male names, I for female names. Note that E now holds family AND male — not the stereotyped pairing.



The Implicit Association Test

	Block 1	Block 2	Block 3
Task description	Initial attribute discrimination	Associated target discrimination	Initial combined task
Example trial	Family Career Wedding	Male Female Daniel	Family Career or or Male Female Salary
Key	● ● E I	● ● E I	● ● E I

Block 3

Four categories, two keys. E if family OR male; I if career OR female. Watch what your hand does.



The Implicit Association Test

	Block 1	Block 2	Block 3	Block 4
Task description	Initial attribute discrimination	Associated target discrimination	Initial combined task	Reversed attribute discrimination
Example trial	Family Career Wedding	<u>Male</u> Female Daniel	Family Career or or Male Female Salary	<u>Career</u> Family Wedding
Key	● ● E I	● ● E I	● ● E I	● ● E I

Block 4

Block 1's mapping, flipped. Measures nothing — it exists to break the habit before the next block.



The Implicit Association Test

	Block 1	Block 2	Block 3	Block 4	Block 5
Task description	Initial attribute discrimination	Associated target discrimination	Initial combined task	Reversed attribute discrimination	Reversed combined task
Example trial	Family Career Wedding	Male Female Daniel	Family Career <i>or</i> <i>or</i> Male Female Salary	Career Family Wedding	Career Family <i>or</i> <i>or</i> Male Female Salary
Key	  E I	  E I	  E I	  E I	  E I

Block 5

E if career OR male; I if family OR female. For most people this block runs noticeably faster.



What the Test Measures

	Block 1	Block 2	Block 3	Block 4	Block 5
Task description	Initial attribute discrimination	Associated target discrimination	Initial combined task	Reversed attribute discrimination	Reversed combined task
Example trial	Family Career Wedding	Male Female Daniel	Family Career <i>or</i> <i>or</i> Male Female Salary	Career Family Wedding	Career Family <i>or</i> <i>or</i> Male Female Salary
Key	  E I	  E I	  E I	  E I	  E I

$d = (\text{mean latency block 3} - \text{block 5}) / \text{pooled SD}$

After Greenwald, McGhee & Schwartz (1998); figure convention: PsyToolkit.

Difference = IAT effect



How Would You Run This Test on a Model?

The instrument has four parts. Three of them survive unchanged.

A model has no reaction time.



How Would You Run This Test on a Model?

The instrument has four parts. Three of them survive unchanged.

A model has no reaction time.

- Two target word sets — do those carry over?

Ben · Paul · Daniel · John · Jeffrey

Rebecca · Michelle · Emily · Julia · Anna



How Would You Run This Test on a Model?

The instrument has four parts. Three of them survive unchanged.

A model has no reaction time.

Male / Female Names


- Two target word sets — do those carry over?

Career · Corporation · Salary · Office · Professional

- Two attribute word sets — do those carry over?

Wedding · Marriage · Parents · Relatives · Home

✓ Career → Family



How Would You Run This Test on a Model?

The instrument has four parts. Three of them survive unchanged.

A model has no reaction time.

- Two target word sets — do those carry over?
- Two attribute word sets — do those carry over?
- A standardised effect size and a significance test — do those carry over?



How Would You Run This Test on a Model?

The instrument has four parts. Three of them survive unchanged.

A model has no reaction time.

- Two target word sets — do those carry over?
- Two attribute word sets — do those carry over?
- A standardised effect size and a significance test — do those carry over?
- Reaction time — what replaces it?



How Would You Run This Test on a Model?

The instrument has four parts. Three of them survive unchanged.

A model has no reaction time.

- Two target word sets — do those carry over?
- Two attribute word sets — do those carry over?
- A standardised effect size and a significance test — do those carry over?
- Reaction time — what replaces it?



The Port: Reaction Time → Cosine Similarity

Everything else in the instrument is preserved exactly.

Human IAT

Reaction time to a pairing

WEAT

Cosine similarity between word vectors



The Port: Reaction Time → Cosine Similarity

Everything else in the instrument is preserved exactly.

Human IAT

Reaction time to a pairing

Two target categories

WEAT

Cosine similarity between word vectors

Two target word sets X, Y



The Port: Reaction Time → Cosine Similarity

Everything else in the instrument is preserved exactly.

Human IAT

Reaction time to a pairing

Two target categories

Two attribute poles

WEAT

Cosine similarity between word vectors

Two target word sets X, Y

Two attribute word sets A, B



The Port: Reaction Time → Cosine Similarity

Everything else in the instrument is preserved exactly.

Human IAT

Reaction time to a pairing

Two target categories

Two attribute poles

Standardised effect size d

WEAT

Cosine similarity between word vectors

Two target word sets X, Y

Two attribute word sets A, B

Standardised effect size d — same shape



The Port: Reaction Time $\rightarrow$ Cosine Similarity

Everything else in the instrument is preserved exactly.

Human IAT

Reaction time to a pairing

Two target categories

Two attribute poles

Standardised effect size d

Significance across subjects

WEAT

Cosine similarity between word vectors

Two target word sets X, Y

Two attribute word sets A, B

Standardised effect size d — same shape

Permutation test over partitions of $X \cup Y$



The Port: Reaction Time → Cosine Similarity

Everything else in the instrument is preserved exactly.

Human IAT

Reaction time to a pairing

Two target categories

Two attribute poles

Standardised effect size d

Significance across subjects

WEAT

Cosine similarity between word vectors

Two target word sets X, Y

Two attribute word sets A, B


Standardised effect size d — same shape

Permutation test over partitions of $X \cup Y$

One substitution; everything else is preserved.



The Space, and What “Close” Means



Drawn in 2-D; the real space has 300 dimensions.

Cosine is an angle, not a distance

Every word is a vector from the origin. $\cos(\mathbf{u}, \mathbf{v})$ depends only on the angle between the two directions — not on how far out the points sit.

So on this picture, “close” means “small angle at the origin.”


The two attribute poles

A career, salary, office, business, professional

B home, parents, children, family, wedding



One Word's Score — $s(w, A, B)$



◆ A career words

◆ B family words


● X male names

● Y female names

$$s(w, A, B) = \text{avg cos}(w, a) - \text{avg cos}(w, b)$$

~~$a \in A$~~ ~~$b \in B$~~ $a \in A$ $b \in B$

How much more w leans toward pole A than toward pole B. One number per word.




θA small $\rightarrow$ cos high; θB large $\rightarrow$ cos low.

$$s(\text{John}, A, B) = 0.31 - 0.06 = +0.25$$



From Scores to an Effect Size



◆ A career words


◆ B family words

● X male names

● Y female names

Number-line scores are schematic ($d \approx 1.84$ here); the reported value is the paper's.

every $s(w, A, B)$, $w \in X \cup Y$



separation between the two means

$$d = \frac{\text{avg } s(x, A, B) - \text{avg } s(y, A, B)}{\text{sd of } s(w, A, B), w \in X \cup Y}$$

On the real word lists: $d = 1.81$, $p < 10^{-3}$



Standard Embeddings, Standard Web Text

WEAT

Primary

- GloVe embeddings (Pennington et al., 2014)
- Common Crawl — about 840 billion tokens
- ≈ 2.2 million distinct tokens, 300 dimensions each



Standard Embeddings, Standard Web Text

Primary

- GloVe embeddings (Pennington et al., 2014)
- Common Crawl — about 840 billion tokens
- ≈ 2.2 million distinct tokens, 300 dimensions each

Replication

- word2vec trained on Google News
- Occupation-gender associations agree across the two: $\rho = 0.88$
- So: not an artefact of one algorithm or one corpus



Standard Embeddings, Standard Web Text

Primary

- GloVe embeddings (Pennington et al., 2014)
- Common Crawl — about 840 billion tokens
- ≈ 2.2 million distinct tokens, 300 dimensions each

Replication

- word2vec trained on Google News
- Occupation-gender associations agree across the two: $\rho = 0.88$
- So: not an artefact of one algorithm or one corpus

Nothing here is special, curated, or adversarial. Standard embeddings, standard web text.



Does It Recover What Nobody Disputes?

Start with associations nobody disputes. If WEAT cannot recover these, it measures nothing.

Target sets	Attributes	d	p
Flowers vs Insects	Pleasant / Unpleasant	1.50	$< 10^{-7}$



Does It Recover What Nobody Disputes?

Start with associations nobody disputes. If WEAT cannot recover these, it measures nothing.

0.8

Target sets	Attributes	d	p
Flowers vs Insects	Pleasant / Unpleasant	<u>1.50</u>	$< 10^{-7}$
Instruments vs Weapons	Pleasant / Unpleasant	1.53	$< 10^{-7}$



Does It Recover What Nobody Disputes?

Start with associations nobody disputes. If WEAT cannot recover these, it measures nothing.

Target sets	Attributes	d	p
Flowers vs Insects	Pleasant / Unpleasant	1.50	$< 10^{-7}$
Instruments vs Weapons	Pleasant / Unpleasant	1.53	$< 10^{-7}$

In social psychology, $d = 0.8$ is a large effect. Both of these are roughly double that.



Does It Recover What Nobody Disputes?

Start with associations nobody disputes. If WEAT cannot recover these, it measures nothing.

Target sets	Attributes	d	p
Flowers vs Insects	Pleasant / Unpleasant	1.50	$< 10^{-7}$
Instruments vs Weapons	Pleasant / Unpleasant	1.53	$< 10^{-7}$

In social psychology, $d = 0.8$ is a large effect. Both of these are roughly double that.

Both effects are large and in the expected direction: the method recovers known associations.



Same Method, Different Word Lists

Identical method. Identical statistics. Different word lists.

Target sets	Attributes	d	p
<u>European-American</u> vs <u>African-American</u> names	Pleasant / Unpleasant	1.41	< 10 ⁻⁸



Are Emily and Greg More Employable than Lakisha and Jamal? A Field Experiment on Labor Market Discrimination

Marianne Bertrand & Sendhil Mullainathan

WORKING PAPER 9873

DOI 10.3386/w9873

ISSUE DATE July 2003

We perform a field experiment to measure racial discrimination in the labor market. We respond with fictitious resumes to help-wanted ads in Boston and Chicago newspapers. To manipulate perception of race, each resume is assigned either a very African American sounding name or a very White sounding name. The results show significant discrimination against African-American names: White names receive 50 percent more callbacks for interviews. We also find that race affects the benefits of a better resume. For



Same Method, Different Word Lists

Identical method. Identical statistics. Different word lists.

Target sets	Attributes	d	p
European-American vs African-American names	Pleasant / Unpleasant	1.41	$< 10^{-8}$
Same, <u>Bertrand–Mullainathan</u> name set	Pleasant / Unpleasant	1.50	$< 10^{-4}$
Same, second attribute set	Pleasant* / Unpleasant*	1.28	$< 10^{-3}$



Same Method, Different Word Lists

Identical method. Identical statistics. Different word lists.

Target sets	Attributes	d	p
European-American vs African-American names	Pleasant / Unpleasant	1.41	$< 10^{-8}$
Same, Bertrand–Mullainathan name set	Pleasant / Unpleasant	1.50	$< 10^{-4}$
Same, second attribute set	Pleasant / Unpleasant	1.28	$< 10^{-3}$
Male vs Female names	Career / Family	1.81	$< 10^{-3}$
Male vs Female terms	Mathematics / Arts	1.06	10^{-2}
Male vs Female terms	Science / Arts	1.24	10^{-2}



Same Method, Different Word Lists

Identical method. Identical statistics. Different word lists.

Target sets	Attributes	d	p
European-American vs African-American names	Pleasant / Unpleasant	1.41	$< 10^{-8}$
Same, Bertrand–Mullainathan name set	Pleasant / Unpleasant	1.50	$< 10^{-4}$
Same, second attribute set	Pleasant / Unpleasant	1.28	$< 10^{-3}$
Male vs Female names	Career / Family	1.81	$< 10^{-3}$
Male vs Female terms	Mathematics / Arts	1.06	10^{-2}
Male vs Female terms	Science / Arts	1.24	10^{-2}

Rows 2 and 3 use the names from the résumé-audit study where identical CVs with African-American names got ~50% fewer callbacks.

The embedding reproduces the structure a field experiment found in human recruiters.



How Do We Know It Was in the Text?

Every measurement so far lives in the embedding. Nobody has counted anything in a corpus.

What we have: bias in the text $\Rightarrow$ bias in the embedding. The corpus was the model's only input.

What we need: bias in the embedding $\Rightarrow$ bias in the text. The converse does not follow deductively.

Both figures: Caliskan, Bryson & Narayanan (2017).



How Do We Know It Was in the Text?

Every measurement so far lives in the embedding. Nobody has counted anything in a corpus.

What we have: bias in the text $\Rightarrow$ bias in the embedding. The corpus was the model's only input.

What we need: bias in the embedding $\Rightarrow$ bias in the text. The converse does not follow deductively.

So this is an inference to the best explanation. Three things carry it.

Both figures: Caliskan, Bryson & Narayanan (2017).



How Do We Know It Was in the Text?

Every measurement so far lives in the embedding. Nobody has counted anything in a corpus.

What we have: bias in the text $\Rightarrow$ bias in the embedding. The corpus was the model's only input.

What we need: bias in the embedding $\Rightarrow$ bias in the text. The converse does not follow deductively.

So this is an inference to the best explanation. Three things carry it.

1 Elimination

There are only two possible sources of structure in it: the corpus, and the training procedure.

Both figures: Caliskan, Bryson & Narayanan (2017).



How Do We Know It Was in the Text?

Every measurement so far lives in the embedding. Nobody has counted anything in a corpus.

What we have: bias in the text $\Rightarrow$ bias in the embedding. The corpus was the model's only input.

What we need: bias in the embedding $\Rightarrow$ bias in the text. The converse does not follow deductively.

So this is an inference to the best explanation. Three things carry it.

1 Elimination

There are only two possible sources of structure in it: the corpus, and the training procedure.

2 Replication across corpora and algorithms

GloVe/Common Crawl vs word2vec/Google News — different objective, different text — $\rho = 0.88$.

Both figures: Caliskan, Bryson & Narayanan (2017).



How Do We Know It Was in the Text?

Every measurement so far lives in the embedding. Nobody has counted anything in a corpus.

What we have: bias in the text $\Rightarrow$ bias in the embedding. The corpus was the model's only input.

What we need: bias in the embedding $\Rightarrow$ bias in the text. The converse does not follow deductively.

So this is an inference to the best explanation. Three things carry it.

1 Elimination

There are only two possible sources of structure in it: the corpus, and the training procedure.

2 Replication across corpora and algorithms

GloVe/Common Crawl vs word2vec/Google News — different objective, different text — $\rho = 0.88$.


3 External validation

WEFAT predicts BLS occupational shares at $\rho = 0.90$ — a statistic the model never saw.

Both figures: Caliskan, Bryson & Narayanan (2017).



WEFAT — One Word, Not Two Sets



Same geometry. One thing changes.

WEAT asked: do two **sets** of words lean differently toward A and B?


WEFAT asks: how strongly does **one word** lean — and does that number track something true about the world?

Poles are now the gender words themselves. The scored word is an occupation.

The question changes from “is it there?” to “is it right?”



The One Change to the Formula



WEAT

$$s(w, A, B) = \text{avg } \cos(w, a) - \text{avg } \cos(w, b)$$


WEFAT

$$s(w, A, B) = \frac{\text{avg } \cos(w, a) - \text{avg } \cos(w, b)}{\text{sd of } \cos(w, x), x \in A \cup B}$$

Why divide? Some words sit in a dense part of the space and are mildly similar to everything; others are sharply separated. Dividing by the spread of w 's own similarities to the pole words makes one word's score comparable with another's.




Score Every Occupation



Placements are schematic; the paper's word list and scores are in its supplement.

WEFAT score $\rightarrow$ association with the female pole



Now every occupation word has one number on one axis.

The move: plot it against the share of women actually working in that occupation.

If the model were inventing stereotypes, the plot would be noise.



It Is Not Inventing It — It Is Measuring It



Caliskan, Bryson & Narayanan (2017), Figure 1.

Occupation ↔ gender

y WEFAT score of the occupation word

x % of US workers in that occupation who are women
(Bureau of Labor Statistics, 2015)

$$\rho = 0.90 \quad p < 10^{-18}$$

81% of the variance in a real labour statistic, recovered from word co-occurrence alone — with no labour data anywhere in the training corpus.



How Do We Know It Was in the Text?

Every measurement so far lives in the embedding. Nobody has counted anything in a corpus.

What we have: bias in the text $\Rightarrow$ bias in the embedding. The corpus was the model's only input.

What we need: bias in the embedding $\Rightarrow$ bias in the text. The converse does not follow deductively.

So this is an inference to the best explanation. Three things carry it.

1 Elimination

There are only two possible sources of structure in it: the corpus, and the training procedure.

2 Replication across corpora and algorithms

GloVe/Common Crawl vs word2vec/Google News — different objective, different text — $\rho = 0.88$.

3 External validation

WEFAT predicts BLS occupational shares at $\rho = 0.90$ — a statistic the model never saw.

Both figures: Caliskan, Bryson & Narayanan (2017).



How Do We Know It Was in the Text?

Every measurement so far lives in the embedding. Nobody has counted anything in a corpus.

What we have: bias in the text $\Rightarrow$ bias in the embedding. The corpus was the model's only input.

What we need: bias in the embedding $\Rightarrow$ bias in the text. The converse does not follow deductively.

So this is an inference to the best explanation. Three things carry it.

1 Elimination

There are only two possible sources of structure in it: the corpus, and the training procedure.

2 Replication across corpora and algorithms

GloVe/Common Crawl vs word2vec/Google News — different objective, different text — $\rho = 0.88$.

3 External validation

WEFAT predicts BLS occupational shares at $\rho = 0.90$ — a statistic the model never saw.

Origin — not fidelity. Whether the copy is faithful is the next question.

Both figures: Caliskan, Bryson & Narayanan (2017).



Sources of Bias in LLMs



2 · Training objective


The loss rewards fitting co-occurrence — and amplifies it. Nothing in it opposes bias.



How You Measure Bias in a Prediction

Visual semantic role labelling: who is doing what, to whom, with what, and where.

Five images the model labels COOKING. Look at the AGENT row.



Gender bias for one verb

$$b(\text{verb}, \text{man}) = c(\text{verb}, \text{man}) / [c(\text{verb}, \text{man}) + c(\text{verb}, \text{woman})]$$

0.50 = even · above = male-skewed · below = female-skewed · amplified = further from 0.50

Zhao, Wang, Yatskar, Ordonez & Chang, EMNLP 2017 — Figure 1, imSitu vSRL.



And It Comes Out Stronger Than It Went In

Three verbs, in b(verb, man). Training set on the left, model predictions on the right.

cooking

.33

→

.16

already female-skewed → more so

Zhao et al., EMNLP 2017 — imSitu vSRL, development set.



Trustworthy Large Language Models



Yatin Nandwani

And It Comes Out Stronger Than It Went In

Three verbs, in b(verb, man). Training set on the left, model predictions on the right.

cooking

.33

→

.16

already female-skewed → more so

serving

.402

→

.122

barely skewed → heavily female

Zhao et al., EMNLP 2017 — imSitu vSRL, development set.



Trustworthy Large Language Models



Yatin Nandwani

And It Comes Out Stronger Than It Went In

Three verbs, in b(verb, man). Training set on the left, model predictions on the right.

cooking	.33	→	.16	already female-skewed → more so
----------------	------------	---	------------	---------------------------------

serving	.402	→	.122	barely skewed → heavily female
----------------	-------------	---	-------------	--------------------------------

tuning	.878	→	1.00	heavily male → exclusively male
---------------	-------------	---	-------------	---------------------------------

Zhao et al., EMNLP 2017 — *imSitu vSRL*, development set.



And It Comes Out Stronger Than It Went In

Three verbs, in b(verb, man). Training set on the left, model predictions on the right.

cooking	.33	→	.16	already female-skewed → more so
----------------	------------	---	------------	---------------------------------

serving	.402	→	.122	barely skewed → heavily female
----------------	-------------	---	-------------	--------------------------------

tuning	.878	→	1.00	heavily male → exclusively male
---------------	-------------	---	-------------	---------------------------------

0.050

mean bias amplification across verbs

45.75%

of verbs amplified, not merely reproduced

Zhao et al., EMNLP 2017 — imSitu vSRL, development set.



Trustworthy Large Language Models



Yatin Nandwani

And It Comes Out Stronger Than It Went In

Three verbs, in b(verb, man). Training set on the left, model predictions on the right.

cooking	.33	→	.16	already female-skewed → more so
----------------	------------	---	------------	---------------------------------

serving	.402	→	.122	barely skewed → heavily female
----------------	-------------	---	-------------	--------------------------------

tuning	.878	→	1.00	heavily male → exclusively male
---------------	-------------	---	-------------	---------------------------------

0.050 mean bias amplification across verbs

45.75% of verbs amplified, not merely reproduced

The model is more biased than its own training data.

Zhao et al., EMNLP 2017 — *imSitu vSRL*, development set.



Sources of Bias in LLMs



3 · Alignment — SFT / RLHF

A narrow annotator pool. Optimising to an average preference can collapse minority views.



Who Wrote the Demonstrations?

The people whose answers became the target. ~40 contractors hired; 19 answered the survey.

What is your nationality?	
Filipino	22%
Bangladeshi	22%
American	17%
Albanian	5%
Brazilian	5%
Canadian	5%
Colombian	5%
Indian	5%
Uruguayan	5%
Zimbabwean	5%

What is your age?	
18-24	26.3%
25-34	47.4%
35-44	10.5%
45-54	10.5%
55-64	5.3%
65+	0%

Ouyang et al. (2022), Appendix B.3, Table 12 — voluntary anonymous survey.



Who Wrote the Demonstrations?

The people whose answers became the target. ~40 contractors hired; 19 answered the survey.

What is your nationality?

Filipino	22%
Bangladeshi	22%
American	17%
Albanian	5%
Brazilian	5%
Canadian	5%
Colombian	5%
Indian	5%
Uruguayan	5%
Zimbabwean	5%

What is your age?

18-24	26.3%
25-34	47.4%
35-44	10.5%
45-54	10.5%
55-64	5.3%
65+	0%

What the table says

- Two nationalities supply 44% of the pool — Filipino and Bangladeshi, 22% each.
- American is third, at 17%. Seven other nationalities, one respondent each.

Ouyang et al. (2022), Appendix B.3, Table 12 — voluntary anonymous survey.



Who Wrote the Demonstrations?

The people whose answers became the target. ~40 contractors hired; 19 answered the survey.

What is your nationality?

Filipino	22%
Bangladeshi	22%
American	17%
Albanian	5%
Brazilian	5%
Canadian	5%
Colombian	5%
Indian	5%
Uruguayan	5%
Zimbabwean	5%

What is your age?

18-24	26.3%
25-34	47.4%
35-44	10.5%
45-54	10.5%
55-64	5.3%
65+	0%

What the table says

- Two nationalities supply 44% of the pool — Filipino and Bangladeshi, 22% each.
- American is third, at 17%. Seven other nationalities, one respondent each.
- 73.7% are under 35. Nobody is over 64.
- 89.4% hold an undergraduate degree or higher. Nobody holds a doctorate.

Ouyang et al. (2022), Appendix B.3, Table 12 — voluntary anonymous survey.



Who Wrote the Demonstrations?

The people whose answers became the target. ~40 contractors hired; 19 answered the survey.

What is your nationality?

Filipino	22%
Bangladeshi	22%
American	17%
Albanian	5%
Brazilian	5%
Canadian	5%
Colombian	5%
Indian	5%
Uruguayan	5%
Zimbabwean	5%

What is your age?

18-24	26.3%
25-34	47.4%
35-44	10.5%
45-54	10.5%
55-64	5.3%
65+	0%

What the table says

- Two nationalities supply 44% of the pool — Filipino and Bangladeshi, 22% each.
- American is third, at 17%. Seven other nationalities, one respondent each.
- 73.7% are under 35. Nobody is over 64.
- 89.4% hold an undergraduate degree or higher. Nobody holds a doctorate.

Nineteen people. Nobody over 64.

Ouyang et al. (2022), Appendix B.3, Table 12 — voluntary anonymous survey.



Whose Opinions Does the Aligned Model Reflect?

PROMPT

[OPTIONAL CONTEXT W/ PERSONA]

Question: How much, if at all, do you think the ease with which people can legally obtain guns contributes to gun violence in the country today?

- A. A great deal
- B. A fair amount
- C. Not too much
- D. Not at all
- E. Refused

Answer:

LOG PROBS



"A"	-0.6
"B"	-0.8
"C"	-13.4
"D"	-14.8
...	...

Santurkar et al. (2023) — OpinionQA, 1,498 Pew questions, 60 groups, 9 models.



Whose Opinions Does the Aligned Model Reflect?

PROMPT


[OPTIONAL CONTEXT W/ PERSONA]

Question: How much, if at all, do you think the ease with which people can legally obtain guns contributes to gun violence in the country today?

- A. A great deal
- B. A fair amount
- C. Not too much
- D. Not at all
- E. Refused


Answer:

LOG PROBS



"A"	-0.6
"B"	-0.8
"C"	-13.4
"D"	-14.8
...	...

OPINION DISTRIBUTIONS



Santurkar et al. (2023) — OpinionQA, 1,498 Pew questions, 60 groups, 9 models.



Whose Opinions Does the Aligned Model Reflect?

The same survey instrument as a moment ago — pointed at the model instead of at people.

Before alignment

lower income

moderate

Protestant or Roman Catholic

Every base model tested, across two vendors. Same web, same place.

Santurkar et al. (2023) — OpinionQA, 1,498 Pew questions, 60 groups, 9 models.



Whose Opinions Does the Aligned Model Reflect?

The same survey instrument as a moment ago — pointed at the model instead of at people.

Before alignment

lower income

moderate

Protestant or Roman Catholic

Every base model tested, across two vendors. Same web, same place.

After human-feedback tuning

liberal

high income

well educated

not religious

The reported crowdworker profile.

Santurkar et al. (2023) — OpinionQA, 1,498 Pew questions, 60 groups, 9 models.



Whose Opinions Does the Aligned Model Reflect?

The same survey instrument as a moment ago — pointed at the model instead of at people.

Before alignment

lower income

moderate

Protestant or Roman Catholic

Every base model tested, across two vendors. Same web, same place.

After human-feedback tuning

liberal

high income

well educated

not religious

The reported crowdworker profile.

Poorly represented by every model: 65+ · widowed · Mormon

Santurkar et al. (2023) — OpinionQA, 1,498 Pew questions, 60 groups, 9 models.



Whose Opinions Does the Aligned Model Reflect?

The same survey instrument as a moment ago — pointed at the model instead of at people.

Before alignment

lower income

moderate

Protestant or Roman Catholic

Every base model tested, across two vendors. Same web, same place.

After human-feedback tuning

liberal

high income

well educated

not religious

The reported crowdworker profile.

Poorly represented by every model: 65+ · widowed · Mormon


There were no annotators over 64 either.

Santurkar et al. (2023) — OpinionQA, 1,498 Pew questions, 60 groups, 9 models.



Tracing Bias from Pre-Training Data to Alignment

The share of women in each of 220 occupations, measured at five consecutive points.




Thaler et al. (2026), Figure 1 — 220 occupations; U.S. BLS, Dolma, OLMo 7B and its two aligned variants.



Tracing Bias from Pre-Training Data to Alignment

The share of women in each of 220 occupations, measured at five consecutive points.




Thaler et al. (2026), Figure 1 — 220 occupations; U.S. BLS, Dolma, OLMo 7B and its two aligned variants.



Tracing Bias from Pre-Training Data to Alignment

The share of women in each of 220 occupations, measured at five consecutive points.




Thaler et al. (2026), Figure 1 — 220 occupations; U.S. BLS, Dolma, OLMo 7B and its two aligned variants.



Tracing Bias from Pre-Training Data to Alignment

The share of women in each of 220 occupations, measured at five consecutive points.




Thaler et al. (2026), Figure 1 — 220 occupations; U.S. BLS, Dolma, OLMo 7B and its two aligned variants.



Tracing Bias from Pre-Training Data to Alignment

The share of women in each of 220 occupations, measured at five consecutive points.




Thaler et al. (2026), Figure 1 — 220 occupations; U.S. BLS, Dolma, OLMo 7B and its two aligned variants.



Tracing Bias from Pre-Training Data to Alignment

The share of women in each of 220 occupations, measured at five consecutive points.




Thaler et al. (2026), Figure 1 — 220 occupations; U.S. BLS, Dolma, OLMo 7B and its two aligned variants.



Tracing Bias from Pre-Training Data to Alignment

The share of women in each of 220 occupations, measured at five consecutive points.



The same quantity, measured at five stages.

Thaler et al. (2026), Figure 1 — 220 occupations; U.S. BLS, Dolma, OLMo 7B and its two aligned variants.



Women's Representation After Each Alignment Stage

Average difference in women's representation against Dolma, in percentage points.

OLMo 7B base

-8.52

under-represents women, relative to the corpus it was trained on

Thaler et al. (2026), Tables 2–3. Amplification measured against Dolma, not against the previous stage.



Women's Representation After Each Alignment Stage

Average difference in women's representation against Dolma, in percentage points.

OLMo 7B base

-8.52

under-represents women, relative to the corpus it was trained on

+ SFT

+10.53

crosses the line and over-corrects

Thaler et al. (2026), Tables 2–3. Amplification measured against Dolma, not against the previous stage.



Women's Representation After Each Alignment Stage

Average difference in women's representation against Dolma, in percentage points.

OLMo 7B base

-8.52

under-represents women, relative to the corpus it was trained on

+ SFT

+10.53

crosses the line and over-corrects

+ Instruct

+27.12

over-represents women in most occupations

Thaler et al. (2026), Tables 2–3. Amplification measured against Dolma, not against the previous stage.



Women's Representation After Each Alignment Stage

Average difference in women's representation against Dolma, in percentage points.

OLMo 7B base

-8.52

under-represents women, relative to the corpus it was trained on

+ SFT

+10.53

crosses the line and over-corrects

+ Instruct

+27.12

over-represents women in most occupations

Skew never returns to the corpus level. It is inverted and enlarged.

Thaler et al. (2026), Tables 2–3. Amplification measured against Dolma, not against the previous stage.



Women's Representation After Each Alignment Stage

Average difference in women's representation against Dolma, in percentage points.

OLMo 7B base

-8.52

under-represents women, relative to the corpus it was trained on

+ SFT

+10.53

crosses the line and over-corrects

+ Instruct

+27.12

over-represents women in most occupations

Skew never returns to the corpus level. It is inverted and enlarged.

And the ranking survives the correction

Smallest corrections: the male-coded manual trades — installation and repair +7.13, construction +9.95.

Largest: sectors already near balance — life and social science +51.78, architecture and engineering +50.87.

Thaler et al. (2026), Tables 2–3. Amplification measured against Dolma, not against the previous stage.



Sources of Bias in LLMs



4 · Prompt & deployment

System prompt, persona, in-context examples, option order. None of it is in the weights.



The Prompt Changes the Decision

Same model, same weights — 70 decision scenarios, demographics varied in the prompt.

Emotional phrasing, eg. *“I really just want to make the right call here” and “This choice is incredibly important.”* larger measured discrimination

Formal bulleted list smaller measured discrimination

Tamkin et al. (2023), §6–§7. Also: Zheng et al. (2024), option order; Lu et al. (2022), example order; Gupta et al. (2024), persona.



The Pi

Same model


Emotional
*want to
choice is*

Formal

“It is ill
“Ignore

Discrimination Score (Avg. Δ in Logits)

2.0
1.5
1.0
0.5
0.0
-0.5



Tamkin et al. (2021)



Trustworthy Large Language Models



Yatin Nandwani

Sources of Bias in LLMs




5 · Evaluation & feedback

Benchmarks encode their authors' assumptions — and output re-enters tomorrow's corpus.



Output Re-enters the Corpus

The corpus is no longer independent of the models trained on it.




Pew (2026), as earlier. Model collapse: Shumailov et al. (2023). What a benchmark encodes: Wallach et al. (2025).



Output Re-enters the Corpus

The corpus is no longer independent of the models trained on it.




Over one-third of pages published after ChatGPT show signs of AI authorship.

Pew (2026), as earlier. Model collapse: Shumailov et al. (2023). What a benchmark encodes: Wallach et al. (2025).



Output Re-enters the Corpus

The corpus is no longer independent of the models trained on it.



Over one-third of pages published after ChatGPT show signs of AI authorship.


What a model says about a group today is part of what the next model learns about that group.

Pew (2026), as earlier. Model collapse: Shumailov et al. (2023). What a benchmark encodes: Wallach et al. (2025).



Output Re-enters the Corpus

The corpus is no longer independent of the models trained on it.



Over one-third of pages published after ChatGPT show signs of AI authorship.

What a model says about a group today is part of what the next model learns about that group.

Stage 5 closes the loop: today's output is tomorrow's training data.

Pew (2026), as earlier. Model collapse: Shumailov et al. (2023). What a benchmark encodes: Wallach et al. (2025).



What We Did Today

- 1 Four headlines. What exactly is wrong here?

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.
- 4 Fairness is a human construct, and the standard keeps evolving.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.
- 4 Fairness is a human construct, and the standard keeps evolving.
- 5 Societies build systems to ensure it. Algorithms need one too.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.
- 4 Fairness is a human construct, and the standard keeps evolving.
- 5 Societies build systems to ensure it. Algorithms need one too.
- 6 Procedural fairness — and why unawareness is not enough.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.
- 4 Fairness is a human construct, and the standard keeps evolving.
- 5 Societies build systems to ensure it. Algorithms need one too.
- 6 Procedural fairness — and why unawareness is not enough.
- 7 Outcome fairness, individual: similar people, similar decisions.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.
- 4 Fairness is a human construct, and the standard keeps evolving.
- 5 Societies build systems to ensure it. Algorithms need one too.
- 6 Procedural fairness — and why unawareness is not enough.
- 7 Outcome fairness, individual: similar people, similar decisions.
- 8 Outcome fairness, group: demographic parity, equalised odds.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.
- 4 Fairness is a human construct, and the standard keeps evolving.
- 5 Societies build systems to ensure it. Algorithms need one too.
- 6 Procedural fairness — and why unawareness is not enough.
- 7 Outcome fairness, individual: similar people, similar decisions.
- 8 Outcome fairness, group: demographic parity, equalised odds.
- 9 LLMs: one model, many decisions — and text, not only decisions.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.
- 4 Fairness is a human construct, and the standard keeps evolving.
- 5 Societies build systems to ensure it. Algorithms need one too.
- 6 Procedural fairness — and why unawareness is not enough.
- 7 Outcome fairness, individual: similar people, similar decisions.
- 8 Outcome fairness, group: demographic parity, equalised odds.
- 9 LLMs: one model, many decisions — and text, not only decisions.
- 10 Five sources of bias, from the corpus to the feedback loop.

Next lecture opens on this slide.



What We Did Today

- 1 Four headlines. What exactly is wrong here?
- 2 “Bias” means four different things. We meant one of them.
- 3 Bias is a skew that is unfair — so fairness has to come first.
- 4 Fairness is a human construct, and the standard keeps evolving.
- 5 Societies build systems to ensure it. Algorithms need one too.
- 6 Procedural fairness — and why unawareness is not enough.
- 7 Outcome fairness, individual: similar people, similar decisions.
- 8 Outcome fairness, group: demographic parity, equalised odds.
- 9 LLMs: one model, many decisions — and text, not only decisions.
- 10 Five sources of bias, from the corpus to the feedback loop.



Next lecture opens on this slide.

